a Each model is consistent with itself

Krippendorff's α across its 3 runs (darker = more consistent)

	D	R	E	A	D2	S	DREAD-S
LLM A	0.88	0.89	0.90	0.88	0.91	0.91	0.94
LLM B	0.89	0.81	0.82	0.75	0.84	0.93	0.88
LLM C	0.86	0.79	0.84	0.80	0.92	0.93	0.95
LLM D	0.91	0.89	0.93	0.91	0.91	0.91	0.96
LLM E	0.94	0.92	0.93	0.92	0.89	0.94	0.94
LLM F	0.93	0.82	0.84	0.83	0.73	0.94	0.93

b Claude rates likelihood lower, GPT-6 does not

Mean rating, LLM minus published (blue = LLM lower, red = higher)

	D	R	E	A	D2	S
LLM A	-1.2	-0.1	-1.8	-1.1	-0.9	-0.3
LLM B	-1.0	+0.0	-1.4	-0.9	-1.6	-0.4
LLM C	-0.4	+0.2	-1.5	-0.4	-1.2	+0.3
LLM D	-1.0	-0.2	-1.7	-0.8	-1.5	+0.0
LLM E	-1.1	-0.1	-1.7	-1.3	-1.4	-0.1
LLM F	+0.3	+1.5	-0.9	+0.3	-0.0	+0.5

c GPT-6 is closest; none is close

Deviation from the published ratings, five DREAD ratings, sorted

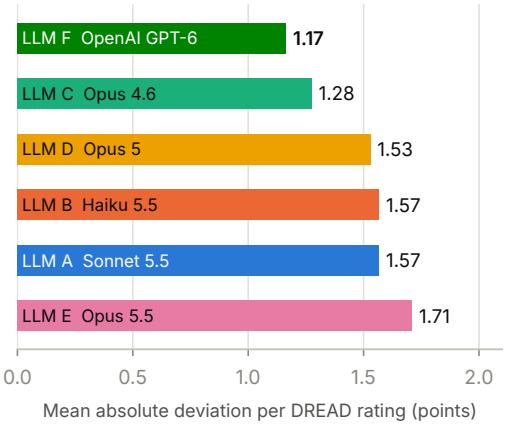

d Claude shifts bands down; GPT-6 stays closer

Distribution of the 21 entries over the DREAD-S bands

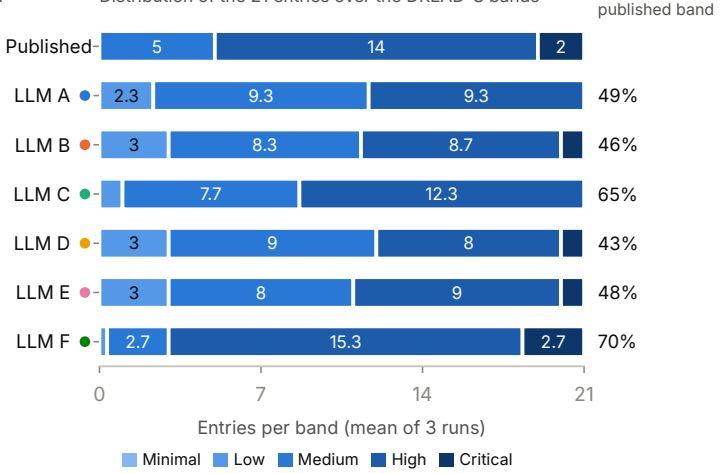

e Ranking only partly kept

τ per run, bar = mean

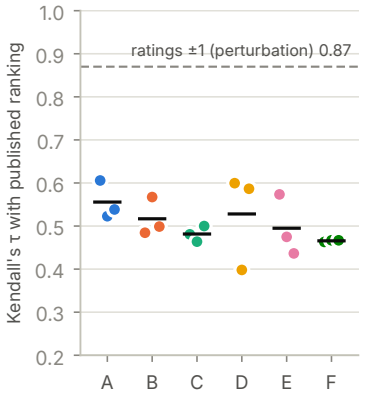

f First place never reproduced

never 1st in any of 18 runs

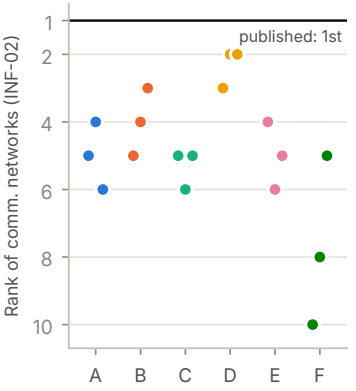

g Safety impact transfers well
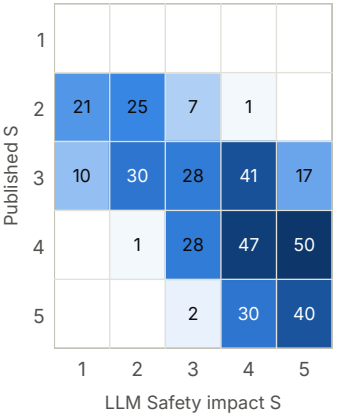
92% of 378 ratings within one level